

Wind Statistics Using Pandas

Introduction:

The data have been modified to contain some missing values, identified by NaN.
Using pandas should make this exercise easier, in particular for the bonus question.

You should be able to perform all of these operations without using a for loop or other looping construct.

1. The data in 'wind.data' has the following format:

In [1]:

```
"""
Yr Mo Dy   RPT   VAL   ROS   KIL   SHA   BIR   DUB   CLA   MUL   CLO   BEL   MAL
61  1  1 15.04 14.96 13.17  9.29   NaN  9.87 13.67 10.25 10.83 12.58 18.50 15.04
61  1  2 14.71   NaN 10.83  6.50 12.62  7.67 11.50 10.04  9.79  9.67 17.54 13.83
61  1  3 18.50 16.88 12.33 10.13 11.17  6.17 11.25   NaN  8.50  7.67 12.75 12.71
"""
```

Out[1]:

```
'\nYr Mo Dy   RPT   VAL   ROS   KIL   SHA   BIR   DUB   CLA   MUL   CLO
BEL   MAL\n61  1  1 15.04 14.96 13.17  9.29   NaN  9.87 13.67 10.25 10.83
12.58 18.50 15.04\n61  1  2 14.71   NaN 10.83  6.50 12.62  7.67 11.50 10.0
4  9.79  9.67 17.54 13.83\n61  1  3 18.50 16.88 12.33 10.13 11.17  6.17 1
1.25   NaN  8.50  7.67 12.75 12.71\n'
```

The first three columns are year, month and day. The remaining 12 columns are average windspeeds in knots at 12 locations in Ireland on that day.

Step 1. Import the necessary libraries

In [1]:

```
import pandas as pd
import numpy as np
import datetime as dt
```

Step 2. Import the dataset

In [33]:

```
pd.read_csv("wind_dataset.csv")
```

Out[33]:

	Yr	Mo	Dy	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAL
0							61	1	1	15.04	14.96	13.17	9.29	NaN	9.87 1...
1							61	1	2	14.71	NaN	10.83	6.50	12.62	7.67 1...
2							61	1	3	18.50	16.88	12.33	10.13	11.17	6.17 1...
3							61	1	4	10.58	6.63	11.75	4.58	4.54	2.88 ...
4							61	1	5	13.33	13.25	11.42	6.17	10.71	8.21 1...
...															...
6569							78	12	27	17.58	16.96	17.62	8.08	13.21	11.67 1...
6570							78	12	28	13.21	5.46	13.46	5.00	8.12	9.42 1...
6571							78	12	29	14.00	10.29	14.42	8.71	9.71	10.54 1...
6572							78	12	30	18.50	14.04	21.29	9.13	12.75	9.71 1...
6573							78	12	31	20.33	17.41	27.29	9.59	12.08	10.13 1...

6574 rows × 1 columns

Step 3. Assign it to a variable called data and replace the first 3 columns by a proper datetime index.

In [46]:

```
data = pd.read_csv("wind_dataset.csv",sep ="\s+",parse_dates= {'Date': ['Yr', 'Mo', 'Dy']})
data
```

Out[46]:

	Date	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAI
0	2061-01-01	15.04	14.96	13.17	9.29	NaN	9.87	13.67	10.25	10.83	12.58	18.50	15.04
1	2061-01-02	14.71	NaN	10.83	6.50	12.62	7.67	11.50	10.04	9.79	9.67	17.54	13.83
2	2061-01-03	18.50	16.88	12.33	10.13	11.17	6.17	11.25	NaN	8.50	7.67	12.75	12.71
3	2061-01-04	10.58	6.63	11.75	4.58	4.54	2.88	8.63	1.79	5.83	5.88	5.46	10.83
4	2061-01-05	13.33	13.25	11.42	6.17	10.71	8.21	11.92	6.54	10.92	10.34	12.92	11.83
...
6569	1978-12-27	17.58	16.96	17.62	8.08	13.21	11.67	14.46	15.59	14.04	14.00	17.21	40.03
6570	1978-12-28	13.21	5.46	13.46	5.00	8.12	9.42	14.33	16.25	15.25	18.05	21.79	41.46
6571	1978-12-29	14.00	10.29	14.42	8.71	9.71	10.54	19.17	12.46	14.50	16.42	18.88	29.53
6572	1978-12-30	18.50	14.04	21.29	9.13	12.75	9.71	18.08	12.87	12.46	12.12	14.67	28.73
6573	1978-12-31	20.33	17.41	27.29	9.59	12.08	10.13	19.25	11.63	11.58	11.38	12.08	22.03

6574 rows × 13 columns



Step 4. Year 2061? Do we really have data from this year? Create a function to fix it and apply it.


In [48]:

```
data["Date"] = np.where(pd.DatetimeIndex(data["Date"]).year < 2000, data.Date, data.Date - data
```

Out[48]:

	Date	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAI
0	1961-01-01	15.04	14.96	13.17	9.29	NaN	9.87	13.67	10.25	10.83	12.58	18.50	15.04
1	1961-01-02	14.71	NaN	10.83	6.50	12.62	7.67	11.50	10.04	9.79	9.67	17.54	13.83
2	1961-01-03	18.50	16.88	12.33	10.13	11.17	6.17	11.25	NaN	8.50	7.67	12.75	12.71
3	1961-01-04	10.58	6.63	11.75	4.58	4.54	2.88	8.63	1.79	5.83	5.88	5.46	10.88
4	1961-01-05	13.33	13.25	11.42	6.17	10.71	8.21	11.92	6.54	10.92	10.34	12.92	11.83
...
6569	1978-12-27	17.58	16.96	17.62	8.08	13.21	11.67	14.46	15.59	14.04	14.00	17.21	40.08
6570	1978-12-28	13.21	5.46	13.46	5.00	8.12	9.42	14.33	16.25	15.25	18.05	21.79	41.46
6571	1978-12-29	14.00	10.29	14.42	8.71	9.71	10.54	19.17	12.46	14.50	16.42	18.88	29.58
6572	1978-12-30	18.50	14.04	21.29	9.13	12.75	9.71	18.08	12.87	12.46	12.12	14.67	28.79
6573	1978-12-31	20.33	17.41	27.29	9.59	12.08	10.13	19.25	11.63	11.58	11.38	12.08	22.08

6574 rows × 13 columns



Step 5. Set the right dates as the index. Pay attention at the data type, it should be datetime64[ns].

In [51]:

```
Data = data.set_index("Date")
Data.index.astype("datetime64[ns]")
```

Out[51]:

```
DatetimeIndex(['1961-01-01', '1961-01-02', '1961-01-03', '1961-01-04',
               '1961-01-05', '1961-01-06', '1961-01-07', '1961-01-08',
               '1961-01-09', '1961-01-10',
               ...,
               '1978-12-22', '1978-12-23', '1978-12-24', '1978-12-25',
               '1978-12-26', '1978-12-27', '1978-12-28', '1978-12-29',
               '1978-12-30', '1978-12-31'],
              dtype='datetime64[ns]', name='Date', length=6574, freq=None)
```

In [52]:

```
Data
```

Out[52]:

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAL
Date												
1961-01-01	15.04	14.96	13.17	9.29	NaN	9.87	13.67	10.25	10.83	12.58	18.50	15.04
1961-01-02	14.71	NaN	10.83	6.50	12.62	7.67	11.50	10.04	9.79	9.67	17.54	13.83
1961-01-03	18.50	16.88	12.33	10.13	11.17	6.17	11.25	NaN	8.50	7.67	12.75	12.71
1961-01-04	10.58	6.63	11.75	4.58	4.54	2.88	8.63	1.79	5.83	5.88	5.46	10.88
1961-01-05	13.33	13.25	11.42	6.17	10.71	8.21	11.92	6.54	10.92	10.34	12.92	11.83
...
1978-12-27	17.58	16.96	17.62	8.08	13.21	11.67	14.46	15.59	14.04	14.00	17.21	40.08
1978-12-28	13.21	5.46	13.46	5.00	8.12	9.42	14.33	16.25	15.25	18.05	21.79	41.46
1978-12-29	14.00	10.29	14.42	8.71	9.71	10.54	19.17	12.46	14.50	16.42	18.88	29.58
1978-12-30	18.50	14.04	21.29	9.13	12.75	9.71	18.08	12.87	12.46	12.12	14.67	28.79
1978-12-31	20.33	17.41	27.29	9.59	12.08	10.13	19.25	11.63	11.58	11.38	12.08	22.08

6574 rows × 12 columns

Step 6. Compute how many values are missing for each location over the entire record.

They should be ignored in all calculations below.

In [64]:

```
Data.isnull().values.ravel().sum()
```

Out[64]:

31

Step 7. Compute how many non-missing values there are in total.

In [65]:

```
x = Data.count()  
print(x.sum())
```

78857

Step 8. Calculate the mean windspeeds of the windspeeds over all the locations and all the times.

A single number for the entire dataset.

In [66]:

```
y = Data.mean()  
y.mean()
```

Out[66]:

10.227982360836924

Step 9. Create a DataFrame called loc_stats and calculate the min, max and mean windspeeds and standard deviations of the windspeeds at each location over all the days

A different set of numbers for each location.

In [67]:

```
def stats(x):
    x = pd.Series(x)
    Min = x.min()
    Max = x.max()
    Mean = x.mean()
    Std = x.std()
    res = [Min,Max,Mean,Std]
    indx = ["Min", "Max", "Mean", "Std"]
    res = pd.Series(res,index =indx)
    return res
loc_stats = Data.apply(stats)
loc_stats
```

Out[67]:

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CL
Min	0.670000	0.210000	1.500000	0.000000	0.130000	0.000000	0.000000	0.000000
Max	35.800000	33.370000	33.840000	28.460000	37.540000	26.160000	30.370000	31.080000
Mean	12.362987	10.644314	11.660526	6.306468	10.455834	7.092254	9.797343	8.495000
Std	5.618413	5.267356	5.008450	3.605811	4.936125	3.968683	4.977555	4.499400

Step 10. Create a DataFrame called day_stats and calculate the min, max and mean windspeed and standard deviations of the windspeeds across all the locations at each day.

A different set of numbers for each day.

In [68]:

```
day_stats=Data.apply(stats, axis=1)
day_stats.head()
```

Out[68]:

	Min	Max	Mean	Std
Date				
1961-01-01	9.29	18.50	13.018182	2.808875
1961-01-02	6.50	17.54	11.336364	3.188994
1961-01-03	6.17	18.50	11.641818	3.681912
1961-01-04	1.79	11.75	6.619167	3.198126
1961-01-05	6.17	13.33	10.630000	2.445356

Step 11. Find the average windspeed in January for each location.

Treat January 1961 and January 1962 both as January.

In [69]:

```
january_data = Data[Data.index.month == 1]
print("January Windspeeds : ")
print(january_data.mean())
```

January Windspeeds :

RPT	14.847325
VAL	12.914560
ROS	13.299624
KIL	7.199498
SHA	11.667734
BIR	8.054839
DUB	11.819355
CLA	9.512047
MUL	9.543208
CLO	10.053566
BEL	14.550520
MAL	18.028763

dtype: float64

Step 12. Downsample the record to a yearly frequency for each location.

In [70]:

```
print("Yearly:\n",Data.resample('A').mean())
```

Yearly:

	RPT	VAL	ROS	KIL	SHA	BI
R \						
Date						
1961-12-31	12.299583	10.351796	11.362369	6.958227	10.881763	7.729726
1962-12-31	12.246923	10.110438	11.732712	6.960440	10.657918	7.393068
1963-12-31	12.813452	10.836986	12.541151	7.330055	11.724110	8.434712
1964-12-31	12.363661	10.920164	12.104372	6.787787	11.454481	7.570874
1965-12-31	12.451370	11.075534	11.848767	6.858466	11.024795	7.478110
1966-12-31	13.461973	11.557205	12.020630	7.345726	11.805041	7.793671
1967-12-31	12.737151	10.990986	11.739397	7.143425	11.630740	7.368164
1968-12-31	11.835628	10.468197	11.409754	6.477678	10.760765	6.067322
1969-12-31	11.166356	9.723699	10.902000	5.767973	9.873918	6.189973
1970-12-31	12.600329	10.726932	11.730247	6.217178	10.567370	7.609452
1971-12-31	11.273123	9.095178	11.088329	5.241507	9.440329	6.097151
1972-12-31	12.463962	10.561311	12.058333	5.929699	9.430410	6.358825
1973-12-31	11.828466	10.680493	10.680493	5.547863	9.640877	6.548740
1974-12-31	13.643096	11.811781	12.336356	6.427041	11.110986	6.809781
1975-12-31	12.008575	10.293836	11.564712	5.269096	9.190082	5.668521
1976-12-31	11.737842	10.203115	10.761230	5.109426	8.846339	6.311038
1977-12-31	13.099616	11.144493	12.627836	6.073945	10.003836	8.586438
1978-12-31	12.504356	11.044274	11.380000	6.082356	10.167233	7.650658

	DUB	CLA	MUL	CLO	BEL	MA
L						
Date						
1961-12-31	9.733923	8.858788	8.647652	9.835577	13.502795	13.68077
1962-12-31	11.020712	8.793753	8.316822	9.676247	12.930685	14.32395
1963-12-31	11.075699	10.336548	8.903589	10.224438	13.638877	14.99901
1964-12-31	10.259153	9.467350	7.789016	10.207951	13.740546	14.91030
1965-12-31	10.618712	8.879918	7.907425	9.918082	12.964247	15.59164
1966-12-31	10.579808	8.835096	8.514438	9.768959	14.265836	16.30726
1967-12-31	10.652027	9.325616	8.645014	9.547425	14.774548	17.13594
1968-12-31	8.859180	8.255519	7.224945	7.832978	12.808634	15.01748
1969-12-31	8.564493	7.711397	7.924521	7.754384	12.621233	15.76290
1970-12-31	9.609890	8.334630	9.297616	8.289808	13.183644	16.45602
1971-12-31	8.385890	6.757315	7.915370	7.229753	12.208932	15.02523
1972-12-31	9.704508	7.680792	8.357295	7.515273	12.727377	15.02871
1973-12-31	8.482110	7.614274	8.245534	7.812411	12.169699	15.44109
1974-12-31	10.084603	9.896986	9.331753	8.736356	13.252959	16.94767
1975-12-31	8.562603	7.843836	8.797945	7.382822	12.631671	15.30786
1976-12-31	9.149126	7.146202	8.883716	7.883087	12.332377	15.47144
1977-12-31	11.523205	8.378384	9.098192	8.821616	13.459068	16.59084

1978-12-31 9.489342 8.800466 9.089753 8.301699 12.967397 16.771370

Step 13. Downsample the record to a monthly frequency for each location.

In [71]:

```
Data.resample('M').mean()
```

Out[71]:

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA
Date								
1961-01-31	14.841333	11.988333	13.431613	7.736774	11.072759	8.588065	11.184839	9.245333
1961-02-28	16.269286	14.975357	14.441481	9.230741	13.852143	10.937500	11.890714	11.846071
1961-03-31	10.890000	11.296452	10.752903	7.284000	10.509355	8.866774	9.644194	9.829677
1961-04-30	10.722667	9.427667	9.998000	5.830667	8.435000	6.495000	6.925333	7.094667
1961-05-31	9.860968	8.850000	10.818065	5.905333	9.490323	6.574839	7.604000	8.177097
...
1978-08-31	9.645161	8.259355	9.032258	4.502903	7.368065	5.935161	5.650323	5.417742
1978-09-30	10.913667	10.895000	10.635000	5.725000	10.372000	9.278333	10.790333	9.583000
1978-10-31	9.897742	8.670968	9.295806	4.721290	8.525161	6.774194	8.115484	7.337742
1978-11-30	16.151667	14.802667	13.508000	7.317333	11.475000	8.743000	11.492333	9.657333
1978-12-31	16.175484	13.748065	15.635161	7.094839	11.398710	9.241613	12.077419	10.194839

216 rows × 12 columns

Step 14. Downsample the record to a weekly frequency for each location.

In [72]:

```
Data.resample('W').mean()
```

Out[72]:

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CL
Date								
1961-01-01	15.040000	14.960000	13.170000	9.290000	NaN	9.870000	13.670000	10.250000
1961-01-08	13.541429	11.486667	10.487143	6.417143	9.474286	6.435714	11.061429	6.616667
1961-01-15	12.468571	8.967143	11.958571	4.630000	7.351429	5.072857	7.535714	6.820000
1961-01-22	13.204286	9.862857	12.982857	6.328571	8.966667	7.417143	9.257143	7.875714
1961-01-29	19.880000	16.141429	18.225714	12.720000	17.432857	14.828571	15.528571	15.160000
...
1978-12-03	14.934286	11.232857	13.941429	5.565714	10.215714	8.618571	9.642857	7.685714
1978-12-10	20.740000	19.190000	17.034286	9.777143	15.287143	12.774286	14.437143	12.488571
1978-12-17	16.758571	14.692857	14.987143	6.917143	11.397143	7.272857	10.208571	7.967143
1978-12-24	11.155714	8.008571	13.172857	4.004286	7.825714	6.290000	7.798571	8.667143
1978-12-31	14.951429	11.801429	16.035714	6.507143	9.660000	8.620000	13.708571	10.477143

940 rows × 12 columns



Step 15. Calculate the min, max and mean windspeeds and standard deviations of the windspeeds across all locations for each week (assume that the first week starts on January 2 1961) for the first 52 weeks.


In [73]:

```
Data = Data.groupby(lambda x: (x.month, x.year))
Data.mean()
```

Out[73]:

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA
(1, 1961)	14.841333	11.988333	13.431613	7.736774	11.072759	8.588065	11.184839	9.245333
(1, 1962)	14.783871	13.160323	12.591935	7.538065	11.779677	8.720000	14.211935	9.600000
(1, 1963)	14.868387	11.112903	15.121613	6.635806	11.080645	7.835484	12.797419	9.844833
(1, 1964)	12.661290	11.818387	11.741290	6.953548	11.400645	6.865806	9.592903	9.687419
(1, 1965)	15.741613	15.546774	15.274194	8.258387	13.588065	9.251290	13.850968	11.260000
...
(12, 1974)	18.511290	17.805806	14.773871	9.734839	16.944194	10.153871	16.602903	15.034194
(12, 1975)	11.655484	8.686774	11.217742	4.478387	6.628710	4.178065	10.351290	6.176129
(12, 1976)	11.962258	10.086774	10.474516	3.383871	7.645484	6.148387	8.034516	4.500000
(12, 1977)	14.751935	12.744839	13.469677	6.592258	11.247742	9.466774	13.231613	10.703871
(12, 1978)	16.175484	13.748065	15.635161	7.094839	11.398710	9.241613	12.077419	10.194833

216 rows × 12 columns



In []: